

Assignment 9

CS 595: Introduction to Web Science

Fall 2013

Shawn M. Jones

Finished on December 4, 2013

1

Question

1. Create a blog-term matrix. Start by grabbing 100 blogs; include:

`http://f-measure.blogspot.com/`

`http://ws-dl.blogspot.com/`

and grab 98 more as per the method shown in class.

Use the blog title as the identifier for each blog (and row of the matrix). Use the terms from every item/title (RSS) or entry/title (Atom) for the columns of the matrix. The values are the frequency of occurrence. Essentially you are replicating the format of the "blogdata.txt" file included with the PCI book code. Limit the number of terms to the most "popular" (i.e., frequent) 500 terms, this is *after* the criteria on p. 32 (slide 7) has been satisfied.

Answer

The script `fetchFeeds.py` shown in Listing 1 outputs a list of URIs that dereference to Atom feeds.

This script uses a canned URI on line 54 to acquire the next blog URI via the call on line 82. If it failed to acquire this URI, it sleeps for 5 seconds and tries again. If successful, it then dereferences the URI using the function defined on line 12. This function returns the representation of the resource, as well as its URI after all redirects have been followed.

If that is successful, it uses the function defined on line 21 to extract the Atom URI from the blog's HTML. If that is successful, it then tries to dereference the Atom feed URI, again acquiring the representation and the redirect-free version of the atom URI. Then it tests the URI using the function defined on line 33. This function checks to ensure that the blog has at least 25 entries. Attempts to ensure that the blog was English did not work well, as all Blogger pages use a character encoding of UTF-8, and detecting 'English' (maybe using a dictionary?) appeared to be outside the realm of this assignment.

Once the URI gets through all of those wickets, it is saved to a list, with the query string "`?max-results=1000`" appended. This ensures that we get at most 1000 entries from each blog, providing the maximum amount of data

for the following questions. Because all blogs come from Blogger, which has a consistent API, this query string works every time.

```
1  #!/usr/local/bin/python
2
3  import sys
4  import os
5  import os.path
6  import feedparser
7  import urllib2
8  import time
9  import chardet
10 from bs4 import BeautifulSoup
11
12 def dereferenceUri(uri):
13
14     pagehandle = urllib2.urlopen(uri)
15     pagedata = pagehandle.read()
16     derefurl = pagehandle.geturl()
17     pagehandle.close()
18
19     return pagedata, derefurl
20
21 def getAtomFeedUri(html):
22
23     soup = BeautifulSoup(html)
24
25     atomLinks = soup.find_all('link',
26                               attrs = { 'rel' : 'alternate', 'type' : 'application/
27                                         atom+xml' })
28
29     # we assume there is only one atom link
30     atomURI = atomLinks[0].attrs['href']
31
32     return atomURI
33
34 def meetsCriteria(feedText):
35
36     parsedData = feedparser.parse(feedText)
37
38     # assume we're good to go by default (fail optimistic?)
39     goodToGo = True
40
41     sys.stderr.write("blog has " + str(len(parsedData.entries))
42                     + " entries\n")
43
44     if (len(parsedData.entries) < 25):
45         goodToGo = False
```

```

45     #if (chardet.detect(feedText)['encoding'] != 'ascii'):
46     #     sys.stderr.write("blog charset is " + chardet.detect(
47         feedText)['encoding'] +
48         ", likely won't parse well for feed vector\n")
49     #     goodToGo = False
50
51     return goodToGo
52
53 def getNextUri():
54     uri = "http://www.blogger.com/next-blog?navBar=true&blogID
55         =3471633091411211117"
56
57     pagehandle = urllib2.urlopen(uri)
58     nexturi = pagehandle.geturl()
59     pagehandle.close()
60
61     return nexturi
62
63 if __name__ == "__main__":
64
65     feedlist = []
66     feedlist.append("http://f-measure.blogspot.com/feeds/posts/
67         default?max-results=200")
68     feedlist.append("http://ws-dl.blogspot.com/feeds/posts/
69         default?max-results=200")
70
71     while (len(feedlist) <= 100):
72
73         try:
74             # Tried these steps, but realized I would have to
75             # parse JavaScript
76             # * look for the iframe containing the next button
77             #     <iframe name='navbar-iframe'...
78             # iframeURI = getIframeUri(html)
79             # * dereference the link from that iframe
80             # iframeText = getIframeText(iframeURI)
81             # * extract the uri form
82             # <a class="b-link" href="...">Next Blog</a>
83             # uri = getNextUri(iframeText)
84             uri = getNextUri()
85         except urllib2.HTTPError as e:
86             sys.stderr.write("failed to acquire next uri,
87                 delaying 5 seconds\n")
88             time.sleep(5)
89         else:
90             sys.stderr.write("working on URI " + uri + "\n")

```

```

89         try:
90             # dereference the uri and get text
91             html,derefuri = dereferenceUri(uri)
92         except urllib2.HTTPError as e:
93             sys.stderr.write("failed to dereference " + uri
94                             +
95                             ", delaying 5 seconds\n")
96             time.sleep(5)
97         else:
98             try:
99                 # fetch the atom feed URI
100                 feedURI = getAtomFeedUri(html)
101                 sys.stderr.write("acquired feed URI " +
102                                 feedURI + "\n")
103             except IndexError as e:
104                 sys.stderr.write(
105                     "failed to acquire Atom feed from HTML,
106                     delaying 5 seconds\n")
107             else:
108                 try:
109                     # get the atom feed text
110                     feedText,feedURI = dereferenceUri(
111                         feedURI)
112                 except urllib2.HTTPError as e:
113                     sys.stderr.write("failed to dereference
114                                     " + feedURI +
115                                     ", delaying 5 seconds\n")
116                     time.sleep(5)
117                 else:
118                     # if it meets the criteria , save the
119                     # file
120                     if meetsCriteria(feedText):
121                         sys.stderr.write("Saving blog feed "
122                                         + feedURI + "?max-results=1000\n")
123                         feedlist.append(feedURI + "?max-
124                                         results=1000")
125
126                     # be nice to the site
127                     time.sleep(1)
128
129     for feed in feedlist:
130         print feed

```

Listing 1: Python script for fetching valid Atom feeds from Blogger

The script is run like so:

```
./fetchFeeds.py > feedlist.txt
```

Once `feedlist.txt` exists, it can be used by Toby Segaran's `generatefeedvector.py`[1]. I only modified it on line 59 so that it could handle UTF-8 encodings for some of the blogs. It is captured in Listing 6 on page 16.

The script `eliminateWords.py` shown in Listing 2 takes in the `blogdata1.txt` file produced by `generatefeedvector.py` and removes the top N terms from it.

It does this by generating scores for each word by summing all of its frequencies across all blogs together. Once it has those scores, it gets the top n words. It then determines the index of each of those words in the list. Using these indices, it regenerates the format of `blogdata.txt`, only printing out each column if it is an “approved” index.

```
1  #!/usr/local/bin/python
2
3  import sys
4
5  sys.path.insert(0, '../libs')
6
7  import clusters
8  import operator
9
10 import pprint
11
12 def getWordscores(words, data):
13
14     wordscores = {}
15
16     for i in range(len(words)):
17         sys.stderr.write('examining ' + words[i] + '\n')
18
19         for j in range(len(data)):
20
21             if words[i] in wordscores:
22                 wordscores[words[i]] += data[j][i]
23             else:
24                 wordscores[words[i]] = data[j][i]
25
26     return wordscores
27
28 def getTopNWords(wordscores, n):
29
30     topNWords = []
31
32     # thanks Stack Overflow:
33     # http://stackoverflow.com/questions/613183/python-sort-a-
34     dictionary-by-value
35     reversedWordscores = sorted(
```

```

35         wordscores.iteritems(), key=operator.itemgetter(1),
36         reverse=True
37     )
38     for i in range(n):
39         sys.stderr.write(
40             "adding " + reversedWordscores[i][0] + " with a
              score of "
41             + str(reversedWordscores[i][1]) + '\n'
42         )
43         topNWords.append(reversedWordscores[i][0])
44
45     return topNWords
46
47
48 if __name__ == '__main__':
49
50     n = int(sys.argv[1])
51
52     blognames, words, data = clusters.readfile('../q1/blogdata1.
        txt')
53
54     wordscores = getWordscores(words, data)
55
56     topNWords = getTopNWords(wordscores, n)
57
58     indexlist = []
59
60     for word in topNWords:
61         indexlist.append(words.index(word))
62
63     lines = []
64
65     line = []
66     line.append('Blog')
67
68     for i in range(len(words)):
69
70         if i in indexlist:
71             line.append(words[i])
72
73     lines.append(line)
74
75     for i in range(len(blognames)):
76         line = []
77         line.append(blognames[i])
78
79         for j in range(len(words)):
80             if j in indexlist:

```

```
81         line.append(str(int(data[i][j])))
82
83     lines.append(line)
84
85     for line in lines:
86         print "\t".join(line)
```

Listing 2: Python script for eliminating words from blog data

2

Question

2. Create an ASCII and JPEG dendrogram that clusters (i.e., HAC) the most similar blogs (see slides 12 & 13). Include the JPEG in your report and upload the ascii file to github (it will be too unwieldy for inclusion in the report).

Answer

The script `makeDendogram.py`, shown in Listing 3 uses Toby Segaran's *clusters.py* [1] shown on page 18.

```
1  #!/usr/local/bin/python
2
3  # all code here stolen shamelessly from
4  # "Programming Collective Intelligence, Chapter 3"
5
6  import sys
7
8  sys.path.insert(0, '../libs')
9
10 import clusters
11
12 blognames, words, data=clusters.readfile('../q1/blogdata500.txt')
13 clust = clusters.hcluster(data)
14
15 # print ASCII dendrogram
16 clusters.printclust(clust, labels=blognames)
17
18 # save JPEG dendrogram
19 clusters.drawdendrogram(clust, blognames, jpeg='blogclust.jpg')
```

Listing 3: Python script for generating dendograms from the blog data captured in question 1

It is run like so:

```
./makeDendogram.py > ascii-dendogram.txt
```

The `printclust` function on line 16 prints out the the dendrogram, so shell redirection is used to save it. The `drawdendrogram` function on line 19 saves a JPEG of the dendrogram, which can be seen in Figure 1.



Figure 1: Dendrogram produced by the *makeDendogram.py* script

Unfortunately, it is difficult to see, but this dendogram shows that the blogs calculated to be most like *F-Measure* are *CRUCIAL CHANGES* and *The Poet You Never Were*. The blog calculated to be most like *Web Science and Digital Libraries Research Group* is *Providence Public Library Computer Classes*.

3

Question

3. Cluster the blogs using K-Means, using k=5,10,20. (see slide 18). How many iterations were required for each value of k?

Answer

The blog clustering is performed by the script shown in Listing 4, which makes use of Toby Segaran's *clusters.py* [1] on page 18., using the function *kcluster* on lines 15, 19, and 23.

```
1  #!/usr/local/bin/python
2
3  # all code here stolen shamelessly from
4  # "Programming Collective Intelligence, Chapter 3"
5
6  import sys
7
8  sys.path.insert(0, '../libs')
9
10 import clusters
11
12 blognames, words, data = clusters.readfile('../q1/blogdata500.txt')
13
14 print "For k=5"
15 kclust = clusters.kcluster(data, k=5)
16 print
17
18 print "For k=10"
19 kclust = clusters.kcluster(data, k=10)
20 print
21
22 print "For k=20"
23 kclust = clusters.kcluster(data, k=20)
24 print
```

Listing 4: Python script for clustering the blogs using K-means, using k=5, 10, and 20

This script is run like so:

```
./makeClusters.py
```

From its output, we see how many iterations each value of k produces.

```
For k=5
Iteration 0
Iteration 1
Iteration 2
Iteration 3
Iteration 4

For k=10
Iteration 0
Iteration 1
Iteration 2
Iteration 3
Iteration 4
Iteration 5
Iteration 6
Iteration 7

For k=20
Iteration 0
Iteration 1
Iteration 2
Iteration 3
```

Thus, for $k = 5$ we get 5 iterations, for $k = 10$ we get 8 iterations, and for $k = 20$ we get 4 iterations.

4

Question

4. Use MDS to create a JPEG of the blogs similar to slide 29. How many iterations were required?

Answer

The blog space is generated using multidimensional scaling from the script `makeMDS.py` shown in Listing 5, which makes use of Toby Segaran's *clusters.py* [1] on page 18, using the functions *scaledown* on line 14 and *draw2d* on line 16.

```
1  #!/usr/local/bin/python
2
3  # all code here stolen shamelessly from
4  # "Programming Collective Intelligence, Chapter 3"
5
6  import sys
7
8  sys.path.insert(0, '../libs')
9
10 import clusters
11
12 blognames, words, data = clusters.readfile('../q1/blogdata500.txt')
13
14 coords = clusters.scaledown(data)
15
16 clusters.draw2d(coords, blognames, jpeg='blogs2d.jpg')
```

Listing 5: Python script for generating a MDS from the blog data

Again, unfortunately the blog space produced does not fit well on a letter-sized page shown in Figure 2. According to this two-dimensional representation, the blog closest to *F-Measure* is *b i t l e y m u s i c*. The blogs closest to *Web Science and Digital Libraries Research Group* are *MUNDO OCHENTAS* and *Providence Public Library Computer Classes*.



Figure 2: Dendrogram produced by the *makeMDS.py* script

5

Question

5. Re-run question 2, but this time with proper TFIDF calculations instead of the hack discussed on slide 7 (p. 32). Use the same 500 words, but this time replace their frequency count with TFIDF scores as computed in assignment #3. Document the code, techniques, methods, etc. used to generate these TFIDF values. Upload the new data file to github.

Compare and contrast the resulting dendrogram with the dendrogram from question #2.

Note: ideally you would not reuse the same 500 terms and instead come up with TFIDF scores for all the terms and then choose the top 500 from that list, but I'm trying to limit the amount of work necessary.

Answer

Not attempted.

A Segaran's *generatefeedvector.py*

```
1 import feedparser
2 import re
3
4 # Returns title and dictionary of word counts for an RSS feed
5 def getwordcounts(url):
6     # Parse the feed
7     d=feedparser.parse(url)
8     wc={}
9
10    # Loop over all the entries
11    for e in d.entries:
12        if 'summary' in e: summary=e.summary
13        else: summary=e.description
14
15        # Extract a list of words
16        words=getwords(e.title+' '+summary)
17        for word in words:
18            wc.setdefault(word,0)
19            wc[word]+=1
20    return d.feed.title,wc
21
22 def getwords(html):
23     # Remove all the HTML tags
24     txt=re.compile(r'<[^>]+>').sub('',html)
25
26     # Split words by all non-alpha characters
27     words=re.compile(r'[^A-Za-z]+').split(txt)
28
29     # Convert to lowercase
30     return [word.lower() for word in words if word!='']
31
32
33 apcount={}
34 wordcounts={}
35 feedlist=[line for line in file('feedlist.txt')]
36 for feedurl in feedlist:
37     try:
38         title,wc=getwordcounts(feedurl)
39         wordcounts[title]=wc
40         for word,count in wc.items():
41             apcount.setdefault(word,0)
42             if count>1:
43                 apcount[word]+=1
44     except:
45         print 'Failed to parse feed %s' % feedurl
46
```

```

47 wordlist=[]
48 for w,bc in apcount.items():
49     frac=float(bc)/len(feedlist)
50     if frac>0.1 and frac<0.5:
51         wordlist.append(w)
52
53 out=file('blogdata1.txt','w')
54 out.write('Blog')
55 for word in wordlist: out.write('\t%s' % word)
56 out.write('\n')
57 for blog,wc in wordcounts.items():
58     print blog
59     blog = blog.encode('UTF-8')
60     out.write(blog)
61     for word in wordlist:
62         if word in wc: out.write('\t%d' % wc[word])
63         else: out.write('\t0 ')
64     out.write('\n')

```

Listing 6: Segaran’s *generatefeedvector.py*, with small modification added on line 59 to handle UTF-8 encodings

B Segaran's *clusters.py*

```
1 from PIL import Image, ImageDraw
2
3 def readfile(filename):
4     lines=[line for line in file(filename)]
5
6     # First line is the column titles
7     colnames=lines[0].strip().split('\t')[1:]
8     rownames=[]
9     data=[]
10    for line in lines[1:]:
11        p=line.strip().split('\t')
12        # First column in each row is the rowname
13        rownames.append(p[0])
14        # The data for this row is the remainder of the row
15        data.append([float(x) for x in p[1:]])
16    return rownames, colnames, data
17
18
19 from math import sqrt
20
21 def pearson(v1, v2):
22     # Simple sums
23     sum1=sum(v1)
24     sum2=sum(v2)
25
26     # Sums of the squares
27     sum1Sq=sum([pow(v,2) for v in v1])
28     sum2Sq=sum([pow(v,2) for v in v2])
29
30     # Sum of the products
31     pSum=sum([v1[i]*v2[i] for i in range(len(v1))])
32
33     # Calculate r (Pearson score)
34     num=pSum-(sum1*sum2/len(v1))
35     den=sqrt((sum1Sq-pow(sum1,2)/len(v1))*(sum2Sq-pow(sum2,2)/len(v1)))
36     if den==0: return 0
37
38     return 1.0-num/den
39
40 class bicluster:
41     def __init__(self, vec, left=None, right=None, distance=0.0, id=None):
42         self.left=left
43         self.right=right
44         self.vec=vec
```

```

45     self.id=id
46     self.distance=distance
47
48 def hcluster(rows,distance=pearson):
49     distances={}
50     currentclustid=-1
51
52     # Clusters are initially just the rows
53     clust=[bicluster(rows[i],id=i) for i in range(len(rows))]
54
55     while len(clust)>1:
56         lowestpair=(0,1)
57         closest=distance(clust[0].vec,clust[1].vec)
58
59         # loop through every pair looking for the smallest distance
60         for i in range(len(clust)):
61             for j in range(i+1,len(clust)):
62                 # distances is the cache of distance calculations
63                 if (clust[i].id,clust[j].id) not in distances:
64                     distances[(clust[i].id,clust[j].id)]=distance(clust[i]
65                                                                    ].vec,clust[j].vec)
66
67                 d=distances[(clust[i].id,clust[j].id)]
68
69                 if d<closest:
70                     closest=d
71                     lowestpair=(i,j)
72
73     # calculate the average of the two clusters
74     mergevec=[
75         (clust[lowestpair[0]].vec[i]+clust[lowestpair[1]].vec[i])
76         /2.0
77         for i in range(len(clust[0].vec))]
78
79     # create the new cluster
80     newcluster=bicluster(mergevec,left=clust[lowestpair[0]],
81                          right=clust[lowestpair[1]],
82                          distance=closest,id=currentclustid)
83
84     # cluster ids that weren't in the original set are negative
85     currentclustid-=1
86     del clust[lowestpair[1]]
87     del clust[lowestpair[0]]
88     clust.append(newcluster)
89
90     return clust[0]
91
92 def printclust(clust,labels=None,n=0):
93     # indent to make a hierarchy layout

```

```

92     for i in range(n): print ' ',
93     if clust.id<0:
94         # negative id means that this is branch
95         print '- '
96     else:
97         # positive id means that this is an endpoint
98         if labels==None: print clust.id
99         else: print labels[clust.id]
100
101     # now print the right and left branches
102     if clust.left!=None: printclust(clust.left ,labels=labels ,n=n
103         +1)
104     if clust.right!=None: printclust(clust.right ,labels=labels ,n=n
105         +1)
106
107 def getheight(clust):
108     # Is this an endpoint? Then the height is just 1
109     if clust.left==None and clust.right==None: return 1
110
111     # Otherwise the height is the same of the heights of
112     # each branch
113     return getheight(clust.left)+getheight(clust.right)
114
115 def getdepth(clust):
116     # The distance of an endpoint is 0.0
117     if clust.left==None and clust.right==None: return 0
118
119     # The distance of a branch is the greater of its two sides
120     # plus its own distance
121     return max(getdepth(clust.left),getdepth(clust.right))+clust.
122         distance
123
124 def drawdendrogram(clust , labels ,jpeg='clusters.jpg') :
125     # height and width
126     h=getheight(clust)*20
127     w=1200
128     depth=getdepth(clust)
129
130     # width is fixed , so scale distances accordingly
131     scaling=float(w-150)/depth
132
133     # Create a new image with a white background
134     img=Image.new( 'RGB' ,(w,h) ,(255,255,255))
135     draw=ImageDraw.Draw(img)
136
137     draw.line((0,h/2,10,h/2) , fill=(255,0,0))
138
139     # Draw the first node

```

```

138 drawnode(draw, clust, 10, (h/2), scaling, labels)
139 img.save(jpeg, 'JPEG')
140
141 def drawnode(draw, clust, x, y, scaling, labels):
142     if clust.id < 0:
143         h1=getheight(clust.left)*20
144         h2=getheight(clust.right)*20
145         top=y-(h1+h2)/2
146         bottom=y+(h1+h2)/2
147         # Line length
148         ll=clust.distance*scaling
149         # Vertical line from this cluster to children
150         draw.line((x, top+h1/2, x, bottom-h2/2), fill=(255,0,0))
151
152         # Horizontal line to left item
153         draw.line((x, top+h1/2, x+ll, top+h1/2), fill=(255,0,0))
154
155         # Horizontal line to right item
156         draw.line((x, bottom-h2/2, x+ll, bottom-h2/2), fill=(255,0,0))
157
158         # Call the function to draw the left and right nodes
159         drawnode(draw, clust.left, x+ll, top+h1/2, scaling, labels)
160         drawnode(draw, clust.right, x+ll, bottom-h2/2, scaling, labels)
161     else:
162         # If this is an endpoint, draw the item label
163         draw.text((x+5, y-7), labels[clust.id], (0,0,0))
164
165 def rotatematrix(data):
166     newdata=[]
167     for i in range(len(data[0])):
168         newrow=[data[j][i] for j in range(len(data))]
169         newdata.append(newrow)
170     return newdata
171
172 import random
173
174 def kcluster(rows, distance=pearson, k=4):
175     # Determine the minimum and maximum values for each point
176     ranges=[(min([row[i] for row in rows]), max([row[i] for row in
177         rows]))
178     for i in range(len(rows[0]))]
179
180     # Create k randomly placed centroids
181     clusters=[[random.random()*(ranges[i][1]-ranges[i][0])+ranges[
182         i][0]
183     for i in range(len(rows[0]))] for j in range(k)]
184
185     lastmatches=None
186     for t in range(100):

```

```

185     print 'Iteration %d' % t
186     bestmatches=[] for i in range(k)]
187
188     # Find which centroid is the closest for each row
189     for j in range(len(rows)):
190         row=rows[j]
191         bestmatch=0
192         for i in range(k):
193             d=distance(clusters[i],row)
194             if d<distance(clusters[bestmatch],row): bestmatch=i
195             bestmatches[bestmatch].append(j)
196
197     # If the results are the same as last time, this is complete
198     if bestmatches==lastmatches: break
199     lastmatches=bestmatches
200
201     # Move the centroids to the average of their members
202     for i in range(k):
203         avgs=[0.0]*len(rows[0])
204         if len(bestmatches[i])>0:
205             for rowid in bestmatches[i]:
206                 for m in range(len(rows[rowid])):
207                     avgs[m]+=rows[rowid][m]
208             for j in range(len(avgs)):
209                 avgs[j]/=len(bestmatches[i])
210             clusters[i]=avgs
211
212     return bestmatches
213
214 def tanamoto(v1,v2):
215     c1,c2,shr=0,0,0
216
217     for i in range(len(v1)):
218         if v1[i]!=0: c1+=1 # in v1
219         if v2[i]!=0: c2+=1 # in v2
220         if v1[i]!=0 and v2[i]!=0: shr+=1 # in both
221
222     return 1.0-(float(shr)/(c1+c2-shr))
223
224 def scaledown(data,distance=pearson,rate=0.01):
225     n=len(data)
226
227     # The real distances between every pair of items
228     realdist=[[distance(data[i],data[j]) for j in range(n)]
229               for i in range(0,n)]
230
231     # Randomly initialize the starting points of the locations in
232     2D
233     loc=[[random.random(),random.random()] for i in range(n)]

```

```

233 fakedist=[[0.0 for j in range(n)] for i in range(n)]
234
235 lasterror=None
236 for m in range(0,1000):
237     # Find projected distances
238     for i in range(n):
239         for j in range(n):
240             fakedist[i][j]=sqrt(sum([pow(loc[i][x]-loc[j][x],2)
241                                     for x in range(len(loc[i]))]))
242
243     # Move points
244     grad=[[0.0,0.0] for i in range(n)]
245
246     totalerror=0
247     for k in range(n):
248         for j in range(n):
249             if j==k: continue
250             # The error is percent difference between the distances
251             errorterm=(fakedist[j][k]-realdist[j][k])/realdist[j][k]
252
253             # Each point needs to be moved away from or towards the
254             # point in proportion to how much error it has
255             grad[k][0]+=((loc[k][0]-loc[j][0])/fakedist[j][k])*
256                        errorterm
257             grad[k][1]+=((loc[k][1]-loc[j][1])/fakedist[j][k])*
258                        errorterm
259
260             # Keep track of the total error
261             totalerror+=abs(errorterm)
262         print totalerror
263
264     # If the answer got worse by moving the points, we are done
265     if lasterror and lasterror<totalerror: break
266     lasterror=totalerror
267
268     # Move each of the points by the learning rate times the
269     # gradient
270     for k in range(n):
271         loc[k][0]-=rate*grad[k][0]
272         loc[k][1]-=rate*grad[k][1]
273
274     return loc
275
276 def draw2d(data, labels, jpeg='mds2d.jpg'):
277     img=Image.new('RGB',(2000,2000),(255,255,255))
278     draw=ImageDraw.Draw(img)
279     for i in range(len(data)):
280         x=(data[i][0]+0.5)*1000

```



```
278     y=(data[i][1]+0.5)*1000
279     draw.text((x,y),labels[i],(0,0,0))
280     img.save(jpeg,'JPEG')
```

Listing 7: Segaran's *clusters.py*

References

- [1] SEGARAN, T. *Programming Collective Intelligence*, first ed. O'Reilly, O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472, August 2007.
- [2] SEGARAN, T. *Programming Collective Intelligence*, first ed. O'Reilly, 2007.